

# Muhammad Uzair Afridi

+966508841267 | [uzairafrididev@gmail.com](mailto:uzairafrididev@gmail.com) | [linkedin.com/in/uzair-afridi00](https://www.linkedin.com/in/uzair-afridi00) | [github.com/uzairafridi00](https://github.com/uzairafridi00)

## PROFESSIONAL SUMMARY

Curious, builder-minded **LLM & AI R&D Engineer** with 3+ years of hands-on experience researching, experimenting with, and deploying **Large Language Models** in production. Deep practical knowledge of **LLM capabilities and limitations** — including hallucinations, context management, and prompt architecture — with proven expertise in **prompt engineering**, **RAG pipelines**, **model evaluation**, and **agentic AI frameworks**. Passionate about keeping pace with the latest AI research and rapidly translating state-of-the-art techniques into working prototypes and scalable systems. Strong communicator able to bridge complex AI concepts with engineering and product teams in fast-moving, remote environments.

## EDUCATION

### IMSciences

*Bachelor of Science in Computer Science / Artificial Intelligence*

Peshawar, Pakistan

*Sept. 2019 – July 2023*

### Edwards College

*Pre-Engineering*

Peshawar, Pakistan

*Aug. 2017 – July 2019*

## EXPERIENCE

### JMM Technologies

*Machine Learning Engineer*

Remote

*June 2024 – March 2026*

- Researched, tested, and benchmarked state-of-the-art **LLMs** — both open-source and proprietary — evaluating model capabilities, hallucination behavior, context window limitations, and output reliability across diverse production use cases
- Designed and iterated on robust **prompt architectures** and **prompt engineering** strategies, optimizing model outputs for accuracy, consistency, and task-specific performance across **OpenAI GPT**, **Claude**, and open-source models
- Built and rapidly prototyped **RAG pipelines** using **LangChain**, **Pinecone**, and **LLMs**, experimenting with chunking strategies, embedding models, and retrieval techniques to improve retrieval accuracy by over **70%**
- Developed and deployed **agentic AI frameworks** using **LangChain**, **LangGraph**, and **CrewAI** to build multi-step autonomous agents, continuously iterating based on real-world performance and the latest agent research
- Integrated **LLM capabilities** into scalable applications via **FastAPI** and **Docker** on **AWS**, collaborating with engineering teams to ensure seamless system integration and inference optimization for low-latency delivery
- Continuously monitored the AI research landscape, translating cutting-edge techniques from papers and community developments into practical R&D experiments and production-ready improvements
- Contributed to technical documentation, best practices, and architecture decisions — clearly communicating complex AI findings to both technical and non-technical team members

### DevMaze

*ML Intern*

Peshawar, Pakistan

*Nov. 2023 – Jan. 2024*

- Rapidly prototyped and experimented with **LLM-powered applications** using **LangChain**, **OpenAI**, and **Hugging Face**, applying **prompt engineering** and context engineering to optimize multi-step AI task performance
- Evaluated and benchmarked **RAG pipeline** configurations — testing chunking strategies, retrieval methods, and re-ranking approaches — improving retrieval accuracy by **70%** through structured experimentation
- Built and fine-tuned **ML models** for regression and classification using **Python**, **scikit-learn**, and **PyTorch**, improving model performance by up to **25%** through feature engineering and hyperparameter optimization
- Developed reproducible experimentation pipelines and documented research findings, methodologies, and model insights for cross-team knowledge sharing

### Ray Media Group

*Junior Python Developer*

Peshawar, Pakistan

*June 2022 – June 2023*

- Built **RESTful APIs** and **Python**-based backend services, integrating third-party **AI APIs** and automating internal workflows to support client-facing application requirements
- Applied strong **software engineering fundamentals** — code reviews, testing, deployment, and documentation — ensuring system reliability and maintainability across production environments
- Worked with **MySQL** and **PostgreSQL** databases for query optimization and data pipeline design, building foundational skills in scalable system integration

## ACHIEVEMENTS

---

Led **Google Developer Student Club (GDSC)** and **Microsoft Learn Student Ambassador (MLSA)** initiatives — organizing sessions on LLMs, AI research, and ML best practices, and demonstrating the ability to translate complex AI concepts into plain language for diverse audiences.

Delivered hands-on workshops on agentic AI, RAG systems, and LLM applications, consistently staying ahead of the rapidly evolving AI landscape and sharing practical insights with peers and teams.

## TECHNICAL SKILLS

---

**LLMs & R&D:** LLM Research, Model Benchmarking, Prompt Engineering, Prompt Architecture, Fine-Tuning, Hallucination Mitigation, Context Management

**RAG & Vector Databases:** RAG Pipelines, Pinecone, Weaviate (familiar), Chunking Strategies, Embedding Models, Retrieval Optimization

**Agentic Frameworks:** LangChain, LangGraph, CrewAI, PraiseAI, AutoGen, Multi-Agent Orchestration, n8n

**ML & Deep Learning:** PyTorch, scikit-learn, XGBoost, Hugging Face, Pandas, NumPy, Model Evaluation

**Deployment & Integration:** FastAPI, Docker, Kubernetes (basic), REST APIs, CI/CD, Git

**Cloud:** AWS (EC2, Lightsail), GCP, Google Cloud Storage

**Languages:** Python, SQL, Bash

## PERSONAL INFORMATION

---

**Nationality:** Pakistani

**Visa Status:** Valid Iqama

**Notice Period:** Immediate Join